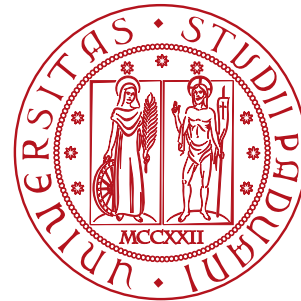


Efficient Low Diameter Clustering

with strong diameter in the CONGEST model

Christian Micheletti



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

1. Distributed Algorithms
2. Distributed Algorithms
3. Network Decomposition
4. Computing a Clustering
5. Computing a Clustering
6. Computing a Clustering

- We want to solve graph problems on **networks**
 - Computers are like nodes in a graph

Distribution \Rightarrow Multiple processors

- We want to solve graph problems on **networks**
 - Computers are like nodes in a graph

Distribution \Rightarrow Multiple processors



Each node gives a partial solution

- At the end they are all combined altogether


...how can we combine those partial solutions?




Distribution \Rightarrow Collaboration

- Arcs are **direct links** between computers
 - Computers have to exchange **messages**

- In the **PN-Network** a node only knows some *Numbered Ports*
 - Connected with a **different** nodes
 - Such nodes are called *neighbours*
- There are no **self loops**




- In the **PN-Network** a node only knows some *Numbered Ports*
 - Connected with a **different** nodes
 - Such nodes are called ***neighbours***
- There are no **self loops**



A node can't see the whole topology

- In the **PN-Network** a node only knows some *Numbered Ports*
 - Connected with a **different** nodes
 - Such nodes are called ***neighbours***
- There are no **self loops**



A node can't see the whole topology



All nodes appear identical



We add **unique identifiers** to the nodes

$$id : V \rightarrow \mathbb{N}$$

where $\forall v \in V : id(v) \leq n^c$ for some $c \geq 1$

We choose n^c so that we need $O(\log n)$ bits to represent an identifier, i.e. identifiers are reasonably “*small*”

- Collaboration requires **exchanging messages**
 - ...on a medium that is **slow** and **unreliable**



⇒ Communication is the main pitfall

- We **measure** the number of messages that an algorithm requires
 - An ***“efficient”*** algorithm will need few messages

W.l.o.g.¹ we adopt a model of ***synchronous communication***

Each round, a node $v \in V$ performs these actions:

1. v ***sends*** a message $msg \in \mathbb{N}$ to its neighbours;
2. v ***receives*** messages from its neighbours;
3. ...

¹Without loss of generality.

W.l.o.g.² we adopt a model of ***synchronous communication***

Each round, a node $v \in V$ performs these actions:

1. v ***sends*** a message $msg \in \mathbb{N}$ to its neighbours;
 2. v ***receives*** messages from its neighbours;
 3. ...
- (1.) and (2.) establish a ***communication round***
 - Main measure of complexity
 - Few communication rounds \Rightarrow few messages

²Without loss of generality.

3. v ***executes locally*** some algorithm (same for each node).
 - A node can ***stop*** in this phase
 - Its local result won't change anymore



(3.) doesn't affect the algorithm's complexity

3. v **executes locally** some algorithm (same for each node).
 - A node can **stop** in this phase
 - Its local result won't change anymore



(3.) doesn't affect the algorithm's complexity

- When all nodes **stopped** the algorithm terminates

We say that an algorithm is “*efficient*” when it stops in **polylogarithmic** number of rounds



- The node with $\text{id}(v) = 1$ “waves *hello*” to neighbours
 - ...sending them a message




- The node with $\text{id}(v) = 1$ “waves hello” to neighbours
 - ...sending them a message
- When a node receives the message, **forwards** it to its neighbours
 - And then **stops**



- The node with $\text{id}(v) = 1$ “waves hello” to neighbours
 - ...sending them a message
- When a node receives the message, **forwards** it to its neighbours
 - And then **stops**
- The running time of this algorithm on a graph G is $O(\text{diam}(G))$

Example: **Maximal Independent Set (MIS)**



- Solving it **centralized** is easy
- How can we solve it **distributed**?



Let's leverage $\text{id}(v)$ to select the next MIS node

- At round $\#i$, node $v : \text{id}(v) = i$ executes
 - If no neighbour is in the MIS, add the node
 - And inform the neighbours
 - Otherwise, the node is outside the MIS



Let's leverage $\text{id}(v)$ to select the next MIS node



- At round $\#i$, node $v : \text{id}(v) = i$ executes
 - If no neighbour is in the MIS, add the node
 - And inform the neighbours
 - Otherwise, the node is outside the MIS
- It is correct since no node has the same id
- This algorithm runs in $O(n^c)$ (the maximum id)
 - **Very bad**



Running a centralized algorithm on a single node would take $O(1)$ rounds

- We'd like to run a MIS algorithm on each node
 - Each must have a **local copy** of the **entire** graph
 - The algorithm must be deterministic
 - When a node stops it checks if it is included in MIS

- The algorithm `GATHER-ALL` makes all nodes build a local copy of the whole graph
 - At round i , each node v knows $ball(i, v)$



- All nodes will know the whole graph after $O(diam(G))$ rounds

- GATHER-ALL assumes that messages size is **unbounded**



Requires to send the whole graph in one message

- It is not always possible to send arbitrary large messages
 - Heavy ones may be “sharded”
- We provide an upper bound for message size
 - Messages need to be reasonably **“small”**
 - Large messages will require more rounds to be sent

In the CONGEST model, messages size has to be $O(\log n)$

- Sending k identifiers takes $O(1)$ rounds
- Sending a set of identifiers can take up to $O(n)$
- Sending the whole graph requires $O(n^2)$ rounds:
 - The adjacency matrix suffices...



⇒ We can't use GATHER-ALL in the CONGEST model

Solving problems in CONGEST


- Censor-Hillel et al. [1] provided an algorithm that solves MIS in $O(\text{diam}(G) \log^2 n)$ in CONGEST



The diameter can be very large

- Worst case: $\text{diam}(G) = n$
- How can we improve it?


- A **Network Decomposition** groups nodes in **colored clusters**
 - Clusters with the same color are not adjacent
 - We say it to **have diameter** d if each cluster has diameter at most d
 - It has c colors






Solving MIS in a color gives a partial solution

- We can apply [1] for all colors
 - (dropping MIS neighbours after each iter)
 - This has complexity $O(c \cdot d \log^2 n)$
 - If $c = O(\log n) = d$ it would be **“efficient”**




A **low diameter clustering** $\mathcal{C} \subseteq 2^V$ for a graph G with diameter d is such:

1. $\forall C_1 \neq C_2 \in \mathcal{C} : \text{dist}_G(C_1, C_2) \geq 2$
 - ***“There are no adjacent clusters”***
2. $\forall C \in \mathcal{C} : \text{diam}(G[C]) \leq d$
 - ***“Any cluster has diameter at most d ”***




Main iteration:

1. Find a low diameter clustering
2. Assign a free color to its nodes
3. Repeat to discarded nodes until there are no more left




Main iteration:

1. Find a low diameter clustering
2. Assign a free color to its nodes
3. Repeat to discarded nodes until there are no more left




Main iteration:

1. Find a low diameter clustering
2. Assign a free color to its nodes
3. Repeat to discarded nodes until there are no more left



Main iteration:

1. Find a low diameter clustering
2. Assign a free color to its nodes
3. Repeat to discarded nodes until there are no more left



- Our previous definition of diameter is also called **strong** diameter

We say a clustering has **weak** diameter when:

1. (unchanged) “There are no adjacent clusters”
2. Any cluster has “**diameter in G** ” at most d

How to compute a low
diameter clustering



- Main accomplishments of [2]:
 - Terminates in $O(\log^6 n)$ rounds in the CONGEST model
 - Outputs a clustering with $O(\log^3 n)$ colors
 - Directly strong diameter



- Main accomplishments of [2]:
 - Terminates in $O(\log^6 n)$ rounds in the CONGEST model
 - Outputs a clustering with $O(\log^3 n)$ colors
 - Directly strong diameter
- Previously [3] provided a l.d.c with **weak** diameter
 - $O(\log^7 n)$ rounds with $O(\log^3 n)$ colors
 - It's possible to turn it into strong diameter




- Main accomplishments of [2]:
 - Terminates in $O(\log^6 n)$ rounds in the CONGEST model
 - Outputs a clustering with $O(\log^3 n)$ colors
 - Directly strong diameter
- Previously [3] provided a l.d.c with **weak** diameter
 - $O(\log^7 n)$ rounds with $O(\log^3 n)$ colors
 - It's possible to turn it into strong diameter
- [4] did it in $O(\log^4 n)$ rounds with $O(\log^3 n)$ colors
 - Has to pass by a weak d. intermediate solution



Objectives:

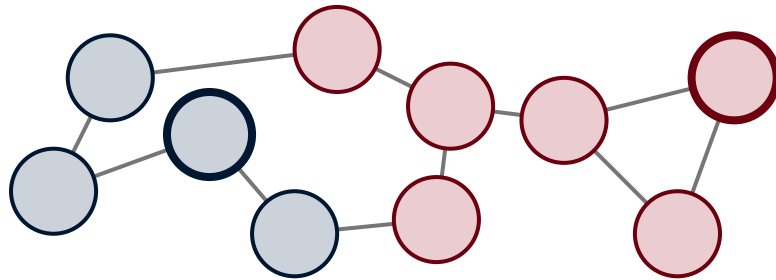
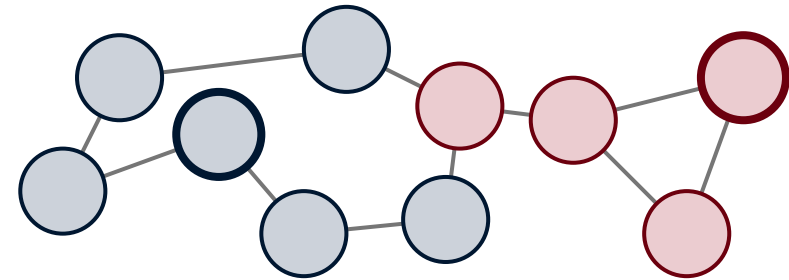
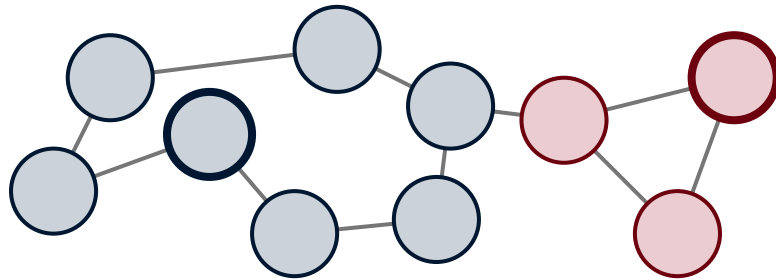
1. Creating connected components with “**low**” diameters
 - Keep track of the “center” of the c.c.
 - A.k.a. **Terminal**
 - We will merge c.c.s
 - Only one terminal is going to be the new center
 - Remove c.c.s with nodes “too far away”





Objectives:

1. Creating connected components with **“low”** diameters
 - Keep track of the “center” of the c.c.
 - A.k.a. **Terminal**
 - We will merge c.c.s
 - Only one terminal is going to be the new center
 - Remove c.c.s with nodes “too far away”
2. Cluster **at least half** of the nodes
 - Required for **“few”** colors **network decompositions**









Phases

- There are $b = \log(\max_{v \in V} \text{id}(v)) = O(\log n)$ **phases**
- "One phase for each bit in index"
 - Phase $i \in [0, b - 1]$ computes **terminals set** Q_i

Notation:

- Q_i is the terminals set built *before* phase i
- Q_b is the terminals set built *after* phase $b - 1$
- V_i is the set of **living nodes** at the beginning of phase i
- $V' = V_b$ is the set of **living nodes** after the last phase



1. Q_i is R_i -ruling, i.e. $\text{dist}_G(Q_i, v) \leq R_i$ for all $v \in V$
 - **We set** $R_i = i * O(\log^2 n)$
 - Q_0 is 0-ruling, trivially true with $Q_0 = V$
 - ***“All nodes are terminals at the beginning”***
 - Q_b is $O(\log^3 n)$ -ruling

“Each node has polylog distance from Q_b ”
 \Rightarrow Each c.c. has at least one terminal



2. Let $q_1, q_2 \in Q_i$ s.t. they are in the same c.c in $G[V_i]$.
Then $\text{id}(q_1)[0..i] = \text{id}(q_2)[0..i]$
- For $i = 0$ it's trivially true
 - For $i = b$ there is ≤ 1 terminal in each c.c.

Along with invariant (1.), it means that each
c.c. has polylog diameter!



3. $|V_i| \geq \left(1 - \frac{i}{2b}\right) |V|$

- $V_0 \geq V$
- $V' \geq \frac{1}{2} |V|$

- ***“The algorithm clusters at least half of the nodes”***

Objective: In a c.c. remove from Q_i all terminals with same $\text{id}(v)$ prefix except one

- Keep c.c. *“small”*
- Divide it if not possible by removing nodes

Outline:

- $2b^2$ *steps*, each computing a forest
- Resulting into a sequence of forests $F_0 \dots F_{2b^2}$

Inductive definition:

- F_0 is a BFS forest with roots set Q_i
- Let T be any tree in F_j and r its root
 - If $\text{id}(r)[i] = 0$ the whole tree is **red**, otherwise **blue**
 - **red** vertexes stay **red**
 - Some **blue** nodes stay **blue**
 - Some others *propose* **red** trees to join
 - If accepted, they become **red**
 - Otherwise, they are **deleted**

Proposal:

$v \in V_j^{propose} \Leftrightarrow v$ is **blue**

$\wedge v$ is the only one in $path(v, root(v))$
that neighbours a **red** node

- Define T_v the (**blue**) subtree rooted at v

v is the only node in T_v that is also in $V_j^{propose}$

Proposal:

- Each node in $V_j^{propose}$ proposes to a **red** neighbour
- Each **red** tree decides to grow or not
 - If it grows, it accepts all proposing subtrees
 - **blue** nodes become **red**
 - If not, all proposing subtrees are **deleted**
- **Criteria:** it decides to grow if gains at least $\frac{|V(T)|}{2b}$ nodes



If a **red** tree doesn't decide to grow, it will
neighbour **red** nodes only

- This means it will be able to delete nodes only once in the whole phase
 - \Rightarrow At most $\frac{|V|}{2b}$ nodes are lost in each phase
 - \Rightarrow After the b phases at most $\frac{|V|}{2}$ nodes are removed

```
1:  $V_0 \leftarrow V$ 
2:  $Q_0 \leftarrow V$ 
3: for  $i \in 0..b - 1$  do
4:   INIT  $F_0$ 
5:   for  $j \in 0..2b^2 - 1$  do
6:     BUILD  $V_j^{propose}$ 
7:      $F_{j+1} \leftarrow \text{STEP}$ 
8:    $V_{i+1} \leftarrow V(F_{2b^2})$ 
9:    $Q_{i+1} \leftarrow \text{roots}(F_{2b^2})$ 
```

$\left. \begin{array}{l} \text{lines 6-7} \end{array} \right\} O(\text{diam}(T_v))$ $\left. \begin{array}{l} \text{lines 5-7} \end{array} \right\} 2b^2 = O(\log^2 n)$ $\left. \begin{array}{l} \text{lines 3-9} \end{array} \right\} b = O(\log n)$

- Recall invariant (1.)
 - $\forall v \in V : \text{dist}_G(Q_i, v) = O(\log^3 n)$, for all $i \in 0..b$
 - Hence, $\text{diam}(T_v) = O(\log^3 n)$, for all $v \in V$
 - Complexity is $\#steps \cdot \#phases \cdot O(\text{diam}(T_v))$
 $= O(\log n) \cdot O(\log^2 n) \cdot O(\log^3 n)$
- \Rightarrow The algorithm runs in $O(\log^6 n)$ communication steps



- We've seen how to build a **low diameter clustering**
 - In $O(\log^6 n)$ communication steps
 - It clusters at least $\frac{n}{2}$ nodes



- We've seen how to build a **low diameter clustering**
 - In $O(\log^6 n)$ communication steps
 - It clusters at least $\frac{n}{2}$ nodes
- We can apply that until all nodes have a color
 - $O(\log n)$ steps and therefore $O(\log n)$ colors
 - **Network decomposition** in $O(\log^7 n)$

- We've seen how to build a **low diameter clustering**
 - In $O(\log^6 n)$ communication steps
 - It clusters at least $\frac{n}{2}$ nodes
- We can apply that until all nodes have a color
 - $O(\log n)$ steps and therefore $O(\log n)$ colors
 - **Network decomposition** in $O(\log^7 n)$
- We solve MIS [1] for each color ($O(\log n) \cdot \dots$)
 - In parallel in the clusters ($\dots \cdot O(\log^3 n \cdot \log^2 n)$)

- We've seen how to build a **low diameter clustering**
 - In $O(\log^6 n)$ communication steps
 - It clusters at least $\frac{n}{2}$ nodes
 - We can apply that until all nodes have a color
 - $O(\log n)$ steps and therefore $O(\log n)$ colors
 - **Network decomposition** in $O(\log^7 n)$
 - We solve MIS [1] for each color ($O(\log n) \cdot \dots$)
 - In parallel in the clusters ($\dots \cdot O(\log^3 n \cdot \log^2 n)$)
- ⇒ We end up solving MIS in $O(\log^7 n)$ rounds

Bibliography

- [1] K. Censor-Hillel, M. Parter, and G. Schwartzman, "Derandomizing Local Distributed Algorithms under Bandwidth Restrictions." [Online]. Available: <https://arxiv.org/abs/1608.01689>
- [2] V. Rozhoň, B. Haeupler, and C. Grunau, "A Simple Deterministic Distributed Low-Diameter Clustering." [Online]. Available: <https://arxiv.org/abs/2210.11784>

- [3] V. Rozhoň and M. Ghaffari, "Polylogarithmic-Time Deterministic Network Decomposition and Distributed Derandomization." [Online]. Available: <https://arxiv.org/abs/1907.10937>
- [4] V. Rozhoň, M. Elkin, C. Grunau, and B. Haeupler, "Deterministic Low-Diameter Decompositions for Weighted Graphs and Distributed and Parallel Applications." [Online]. Available: <https://arxiv.org/abs/2204.08254>